

AN ABSTRACT OF THE THESIS OF

Michael M. Anderson for the degree of Master of Science in Computer Science
presented on June 13, 2013.

Title: Activity Detection on Free-Living Data Using Change-Point Detection

Abstract approved: _____

Weng-Keen Wong

Activity detection on time series sensor data is a rapidly emerging field with a lot of potential real-world application. In previous work researchers have unrealistically assumed that they know at which ticks in a time series the given subject stops doing one activity and starts doing another. In this thesis we explore the feasibility of segmenting realistic free-living time series data using techniques from change-point detection, and then classifying the predicted segments using standard supervised learning techniques. We compare this to the popular approach of splitting the time series into small windows of fixed-length, predicting on each window with a classifier, and then smoothing over the predictions with an HMM. We find that the HMM approach clearly outperforms the change-point detection approach, but that change-point detection may be promising given a modeling assumption that is appropriate to the data.

©Copyright by Michael M. Anderson
June 13, 2013
All Rights Reserved

Activity Detection on Free-Living Data Using Change-Point Detection

by

Michael M. Anderson

A THESIS

submitted to

Oregon State University

in partial fulfillment of
the requirements for the
degree of

Master of Science

Presented June 13, 2013
Commencement June 2013

Master of Science thesis of Michael M. Anderson presented on June 13, 2013.

APPROVED:

Major Professor, representing Computer Science

Director of the School of Electrical Engineering and Computer Science

Dean of the Graduate School

I understand that my thesis will become part of the permanent collection of Oregon State University libraries. My signature below authorizes release of my thesis to any reader upon request.

Michael M. Anderson, Author

ACKNOWLEDGEMENTS

I would like to acknowledge and thank my advisor, Dr. Weng-Keen Wong, for providing me with a clear and overarching vision of this project at all stages of its development, as well as his general expertise and helpfulness with all of the rough patches, sticking points, and unexpected problems that invariably accompany an endeavor of this magnitude. I would also like to thank the other members of my committee, Dr. Alan Fern, Dr. Raviv Raich, and Dr. Hector Vergara, for giving their time to help supervise the final stages of this project. Thanks go to all of my family and friends who encouraged and stood by me along the way. They are too numerous to name but they know who they are.

TABLE OF CONTENTS

	<u>Page</u>
1 Results	1
1.1 Change-Point Detection	1
1.2 HMM	6
1.3 Discussion	10
2 Conclusion	12
Bibliography	12

LIST OF FIGURES

<u>Figure</u>		<u>Page</u>
1.1	OSU Hip Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.	3
1.2	LiME Day 1 Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.	4
1.3	LiME Day 2 Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.	5
1.4	OSU Hip HMM Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.	7
1.5	LiME Day 1 HMM Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.	8
1.6	LiME Day 2 HMM Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.	9

LIST OF FIGURES (Continued)

<u>Figure</u>		<u>Page</u>
1.7	Comparison of change-point detection and HMMs in terms of accuracy and detection time. For a given false positive rate, change-point detection algorithms split a time series into windows of a certain average size, which decreases as the false positive rate increases. From this it is possible to relate false positive rates from the CPD experiments to the window sizes of the HMM experiments. Here is the HMM data with false positive rates of $\{0.033, 0.028, 0.024, 0.021, 0.019, 0.017\}$ that would correspond to average window sizes of $\{10, 12, 14, 16, 18, 20\}$. Both graphs show decision tree results, and both show results from each of the three datasets clustered together over a single false positive rate.	11

Chapter 1: Results

1.1 Change-Point Detection

Results for our change-point detection experiments are given in Figures 1.1-1.3. We hypothesized that the performance of the change-point detection algorithms would depend heavily on the threshold level for change prediction. This was tested by varying the average number of times per second that the algorithms falsely predicted a change. A large number of such false positive rates per second were tested, but for the sake of brevity only a representative sample of $\{0.005, 0.01, 0.05, 0.1\}$ are shown here.

In the OSU Hip experiments, control charts outperformed KLIEP in terms of detection time (Figures 4.1.2, 4.1.4, 4.1.6), while the accuracy results (Figures 4.1.1, 4.1.3, 4.1.5) were mixed. Except when predicted windows are large enough to span across multiple true activities, it is generally expected that accuracy will decrease as false positive rate increases because small windows contain less information and are less discriminative than larger windows. This behavior is seen in the control chart accuracy results (grey bars in Figures 4.1.1, 4.1.3, 4.1.5), but not in the KLIEP accuracy results (white bars in Figures 4.1.1, 4.1.3, 4.1.5). Follow-up experiments showed that KLIEP peaks in accuracy for false positive rates between 0.2 and 0.3 for all three classifiers. KLIEP seemed to perform best on this dataset when it was given many opportunities to predict changes.

Further investigation indicated that across the OSU Hip dataset the KLIEP algorithm was unable to detect many of the different activity changes without a very low score threshold value (and a very high false positive rates). Some qualitative plotting of the OSU Hip data showed that most of its activities have accelerometer amplitude values that strongly resemble draws from a multivariate normal distribution. Since control charts assume that the data is drawn from a distribution that is a member of that family, it is logical that control charts would outperform algorithms with different modeling assumptions on OSU Hip.

In the LiME experiments, KLIEP outperformed control charts in terms of accuracy across the board, and control charts outperformed KLIEP in terms of detection time

across the board. This suggests that in general control charts correctly detected true changes more quickly, but that after a correct change prediction it was more likely to make an incorrect change prediction.

In a few cases (Figures 4.1.2, 4.1.6, 4.2.6) the detection time did not decrease as the false positive rate increased. On the face of it this would seem to be a non-sequitur, but this only happened in cases when accuracy also decreased (Figures 4.1.1, 4.1.5, 4.2.5). Smaller window sizes tend to be correlated with decreased detection times, but it is possible that predicting with smaller windows, if they happen to contain an insufficient amount of discriminative data, can actually increase the time required for the classifier to start correctly predicting the ground-truth activity. Additionally, the given increases in detection time were small and within confidence bounds.

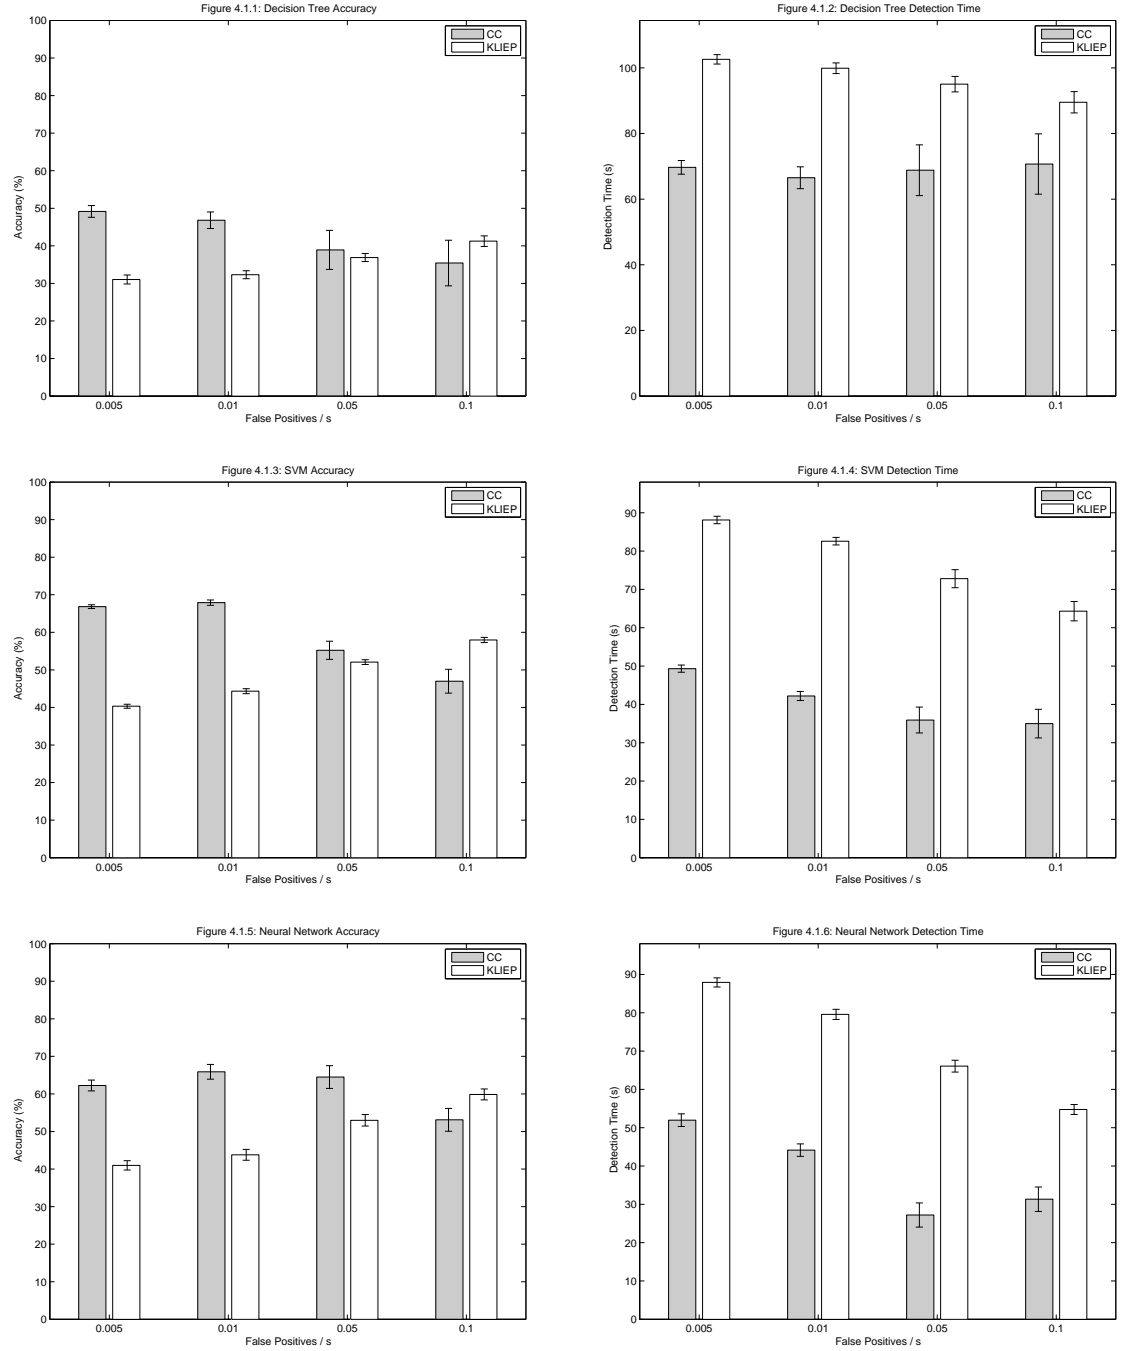


Figure 1.1: OSU Hip Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.




Figure 1.2: LiME Day 1 Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.

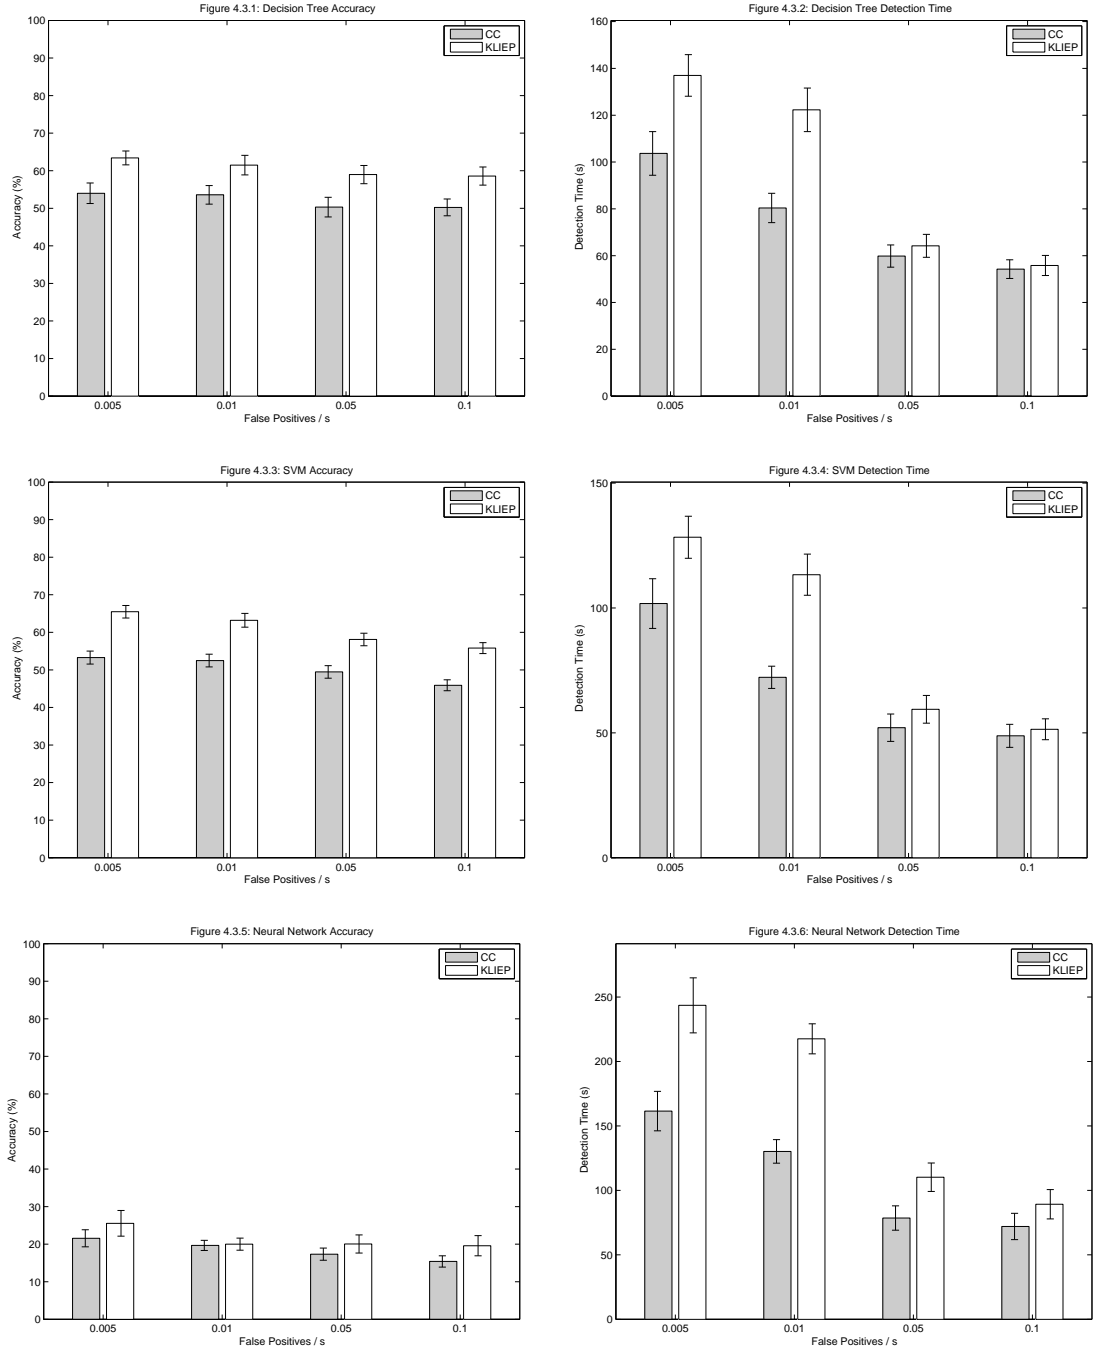


Figure 1.3: LiME Day 2 Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. Change-point detection results were averaged over 30 splits into training, testing, and validation datasets, along with error bars showing a 95% confidence interval.

1.2 HMM

Results for our HMM experiments are given in Figures 1.4-1.6. Each HMM experiment was performed by splitting each time series into windows of fixed length corresponding to discrete time “ticks” in an HMM, and results for windows of length $\{10, 12, 14, 16, 18, 20\}$ seconds are shown.

For both the SVM and decision tree classifiers, accuracy was high and detection time was low across all three datasets. Accuracy and detection time were also stable with respect to window size. Further experiments on the OSU Hip dataset showed that the HMM when paired with these classifiers tends to be stable with window sizes that are greater than a few seconds, which seems to be the amount of time required to be informative. Neural networks performed somewhat more poorly and erratically across the board.

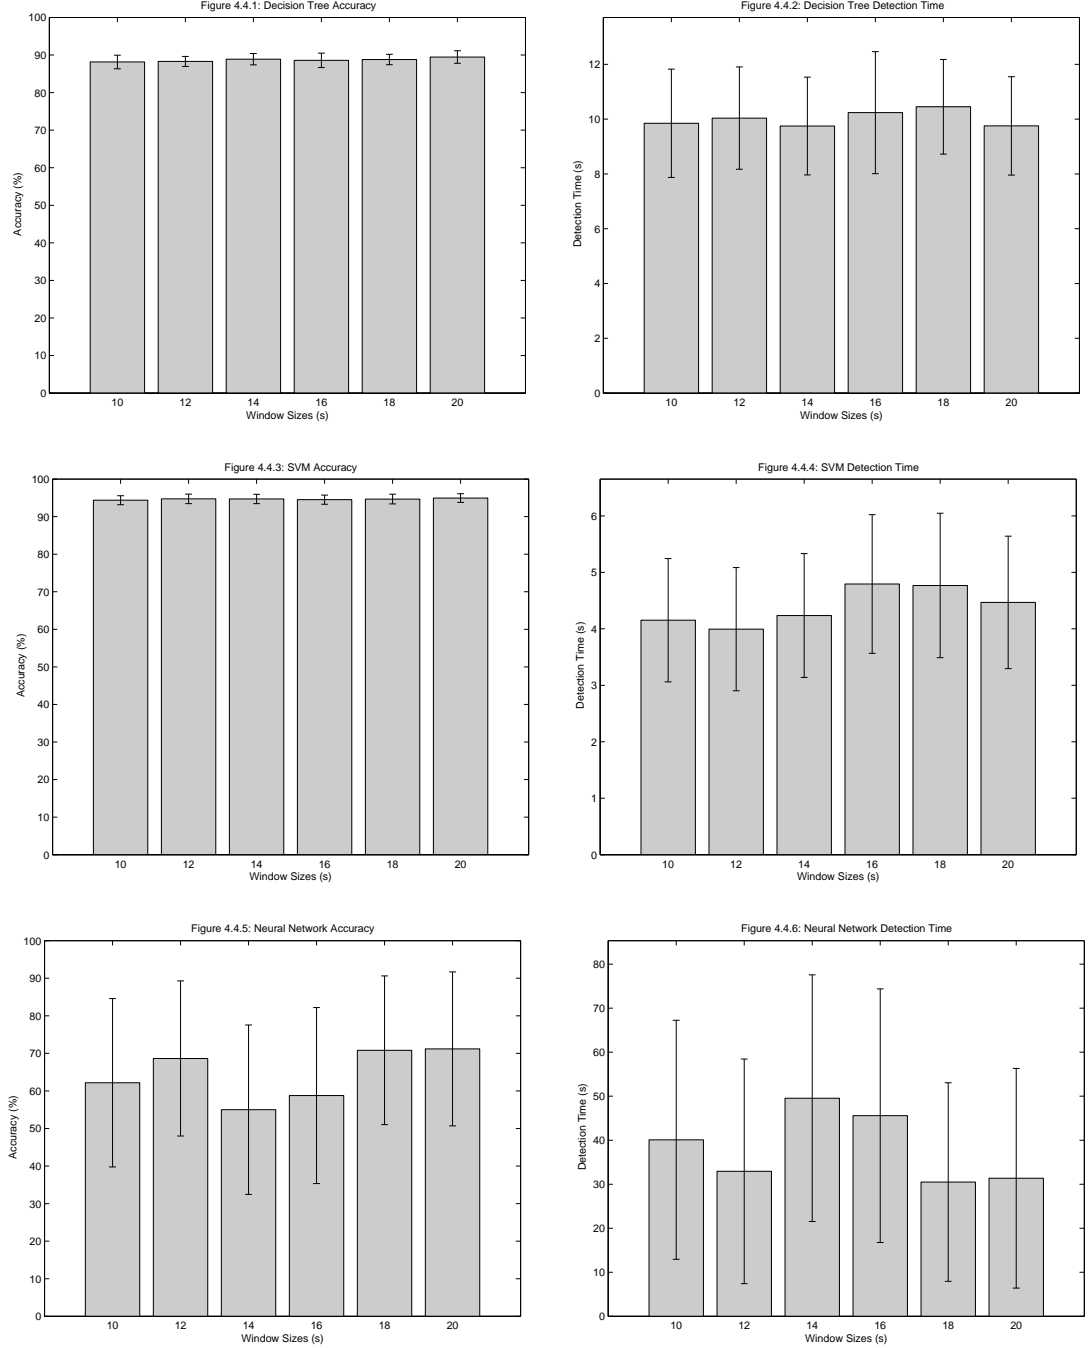


Figure 1.4: OSU Hip HMM Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.

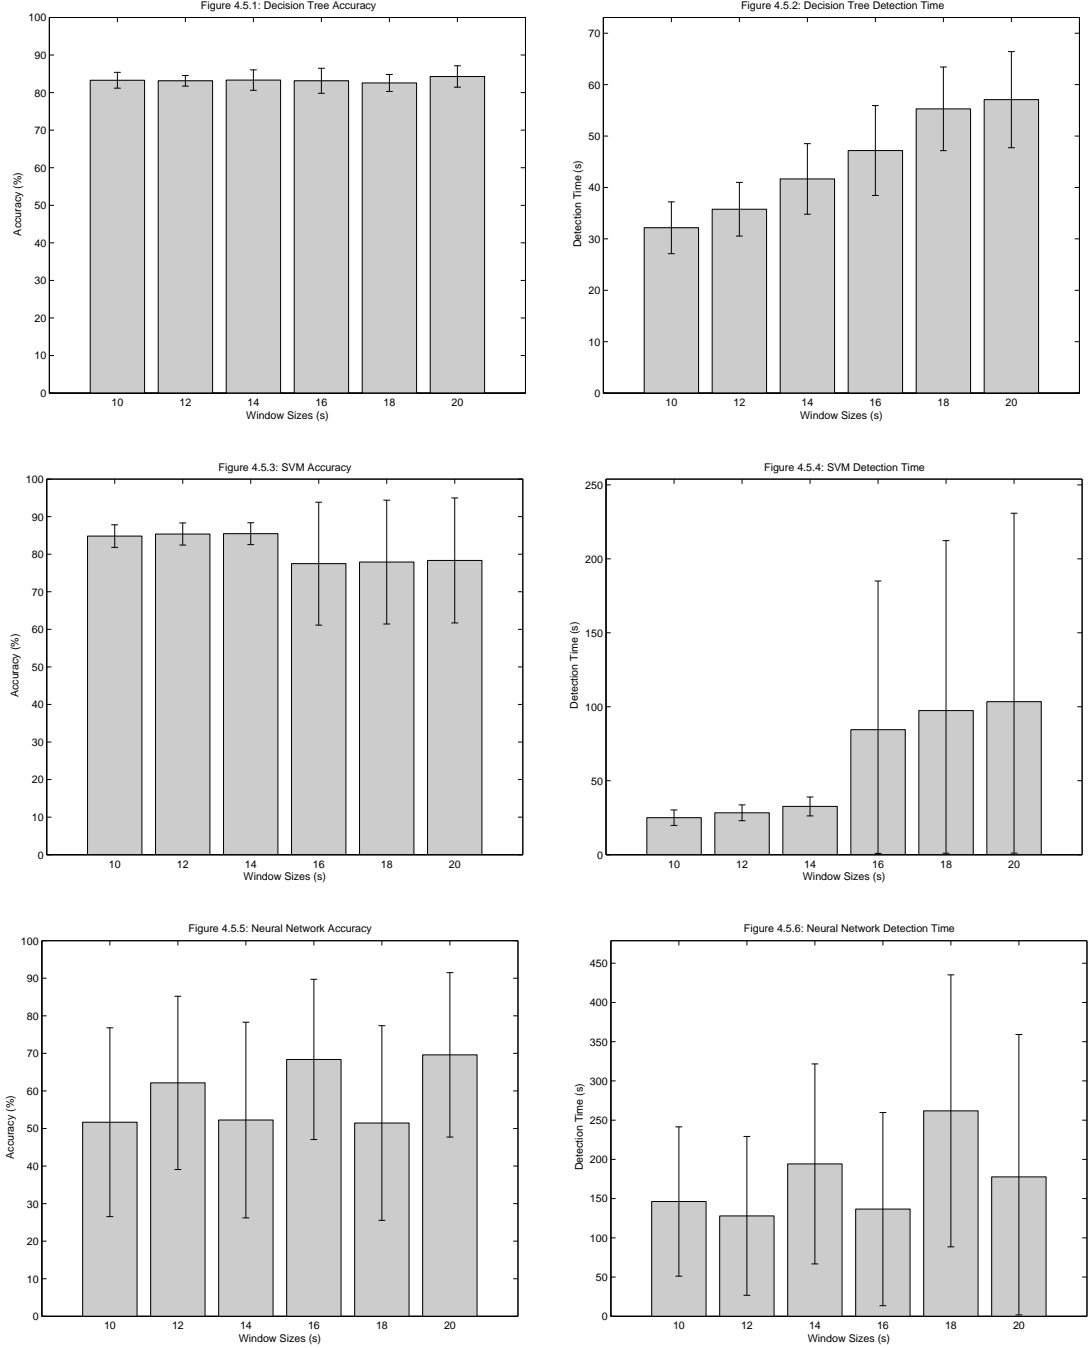


Figure 1.5: LiME Day 1 HMM Results. Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.

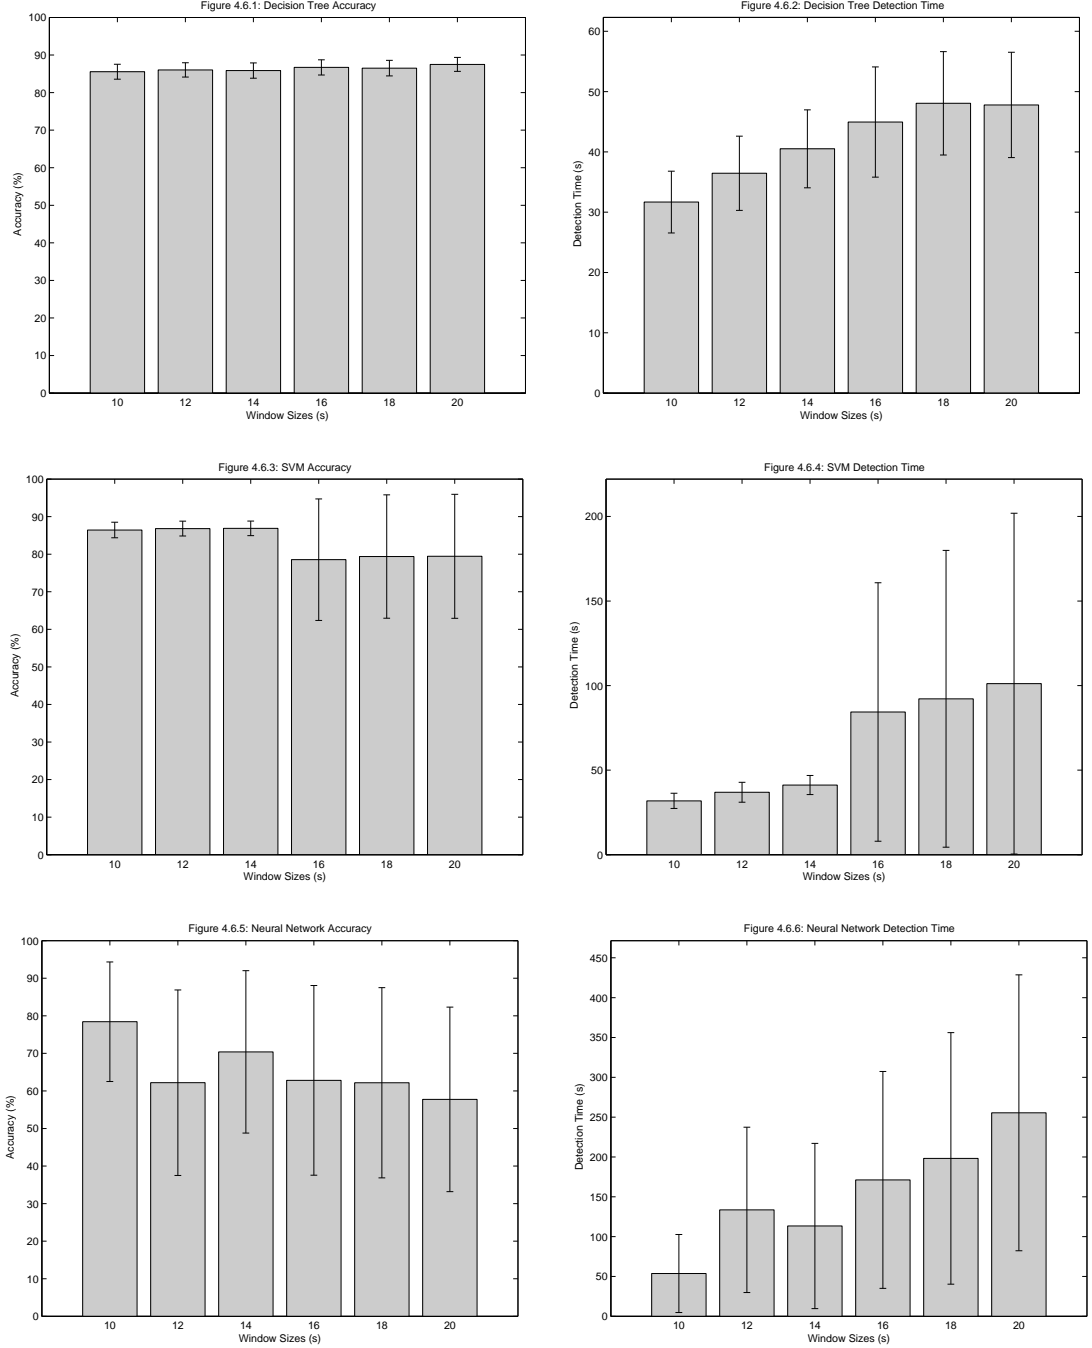


Figure 1.6: LiME Day 2 HMM Results Graphs are organized into rows by base classifier, and columns by evaluation metric. HMM results were averaged over 10 splits into training (base classifier), validation, training (HMM), and testing datasets, along with error bars showing a 95% confidence interval.

1.3 Discussion

Our results clearly show that the HMM approach outperformed the change-point detection approach, both in terms of accuracy and detection time, regardless of the dataset and base classifier. Figure ?? shows a side-by-side comparison of the top-down and bottom-up approaches. HMMs likely did better because they are trained on labelled data. Change-point detection algorithms on the other hand are unaware of true activity labels—they merely signal a change when they detect a distance or dissimilarity between reference and test data. As such, change-point detection algorithms have trouble accurately segmenting free-living data, and noisy segmentation results in poor performance.

A contributing factor to the particularly high accuracy and low detection time results attained for the OSU Hip experiments was that the data consisted of activities that were synthetically glued together. The same group of activities were performed in the same order by each of the 50 subjects in this dataset, making transitions from one activity to the other very predictable for a temporal model. By contrast, the LiME datasets consisted of unsynthetic data gathered from a large set of unstructured and variable-length activities, so the activity transitions were not as predictable and are more indicative of an application of our techniques in the real world.

A final point of interest was that SVM clearly outperformed the other two base classifiers, and that the faster and simpler decision tree model did fairly well against neural networks. This result is significant because much of the previous research that has formulated activity detection as a supervised learning problem has used neural networks exclusively.




Figure 1.7: Comparison of change-point detection and HMMs in terms of accuracy and detection time. For a given false positive rate, change-point detection algorithms split a time series into windows of a certain average size, which decreases as the false positive rate increases. From this it is possible to relate false positive rates from the CPD experiments to the window sizes of the HMM experiments. Here is the HMM data with false positive rates of $\{0.033, 0.028, 0.024, 0.021, 0.019, 0.017\}$ that would correspond to average window sizes of $\{10, 12, 14, 16, 18, 20\}$. Both graphs show decision tree results, and both show results from each of the three datasets clustered together over a single false positive rate.

Chapter 2: Conclusion

The purpose of this work was to test the feasibility of using change-point detection techniques for deciding when one activity ended within a time series and the next began, and to contrast this technique with an HMM approach. The bottom-up approach clearly outperformed the top-down approach. We also showed that the performance of a change-point detection algorithm was highly dependent on how well the data fit the modeling assumption of the algorithm, so it is plausible that a change-point detection algorithm with a modeling assumption that is in accord with the given data will perform well.

Bibliography

- [1] Y. Zheng, W.K. Wong, X. Guan, and S. Trost. Physical activity recognition from accelerometer data using a multi-scale ensemble method. In *Twenty-Fifth Annual Conference on Innovative Applications of Artificial Intelligence*. IAAA, 2013.

